

# Battle of the Neighbourhoods : The Hunt for Hipsters

## 1. Introduction / Business Problem

My partner is a jewellery designer/maker, with a successful shop in a trendy part of London. She has expressed an interest in opening a second shop, and her usual strategy for this might be to go do some market research. Rather than simply ask the coolest people we know, I'd like to use data science to identify an appropriate location.

The question is this: can we work out which are the hippest areas in London, and the areas that are on the verge of becoming hip? And given that where would be a good place to open a cool jewelry shop?

This approach could apply to anyone interested in opening an independent store or brand in London, and if successful could be applied to any city for which we have the appropriate data.

## 2. Data

We'll break this down into 2 components: quantifying trendiness into clusters, and finding competitive advantage.

To quantify trendiness we'll explore different options, among them identifying whether there are a large number of chain (non-independent) stores in the area, or stores with particular or unusual categories. We'll look at the frequency and density of specific category types such as record stores and coffee shops.

To find competitive advantage, we'll look at details of the specific jewelry shops, including rating and price. Consideration should be given to the number of total shops and restaurants in the area (which would generate more footfall and be a positive influence), as well as other jewellery shops (which would mean more competition and be a negative influence).

The data we'll need includes:

- **Geocoding** - to identify the specific locations
- **Foursquare API** - to identify the qualifying places (jewelry stores, restaurants etc.) and essential details around them

### 2.2. Geocoding

[Geonames.org](https://www.geonames.org/) (<https://www.geonames.org/>) is a fantastic free resource that publishes geocoded names and their latitudes/longitudes many countries around the world. The GeoNames geographical database is available for download free of charge under a creative commons attribution license. The resource includes an explicit accuracy for the latitude/longitude, which will be useful for filtering out bad/inaccurate/duplicate data.

Below is a sample of the resulting data. Note the available Latitude/Longitude, and the Accuracy record.

Out[2]:

	Countrycode	Postalcode	Placename	Adminname1	Admincode1	Adminname2	Admincode2
0	GB	DN14	Goole	England	ENG	East Riding of Yorkshire	116090
1	GB	DN14	Pollington	England	ENG	East Riding of Yorkshire	116090
2	GB	DN14	Faxfleet	England	ENG	East Riding of Yorkshire	116090
3	GB	DN14	Laxton	England	ENG	East Riding of Yorkshire	116090
4	GB	DN14	Old Goole	England	ENG	East Riding of Yorkshire	116090

According to the documentation, 'Accuracy' is defined as below:

```
Accuracy is an integer, the higher the better :
1 : estimated as average from numerically neighbouring postal codes
3 : same postal code, other name
4 : place name from geonames db
6 : postal code area centroid
```

Consequently we'll want to limit ourselves to codes with an Accuracy of greater or equal to 3 (averages of neighboring postal codes isn't useful). And because we're going to be using the Latitude/Longitude, we'll remove duplicate records of that and use only the first name. Postalcode will be tricky to use as it bridges Placenames, but we'll want to keep it around for reference and collapse it into a single column joined by a comma.

Finally, we'll be limiting ourselves to Greater London.

Below is a resulting sample of the data we'll use:


Out[3]:

	Latitude	Longitude	Placename	Postalcode
0	51.3133	0.0343	Biggin Hill	TN16
1	51.3148	-0.1570	Hooley	CR5
2	51.3192	0.0712	Cudham	TN14
3	51.3200	-0.1409	Coulsdon	CR5
4	51.3264	-0.1011	Kenley	CR8

There are 468 Places we will evaluate in London

A review of a map of London with the locations marked at a radius of 1000m exposes the amount of overlapping areas and the high coverage for central London. This should be sufficient for our evaluation.

Out[5]:



## 2.2. Foursquare API

We will retrieve both the places and necessary details from the Foursquare API. We'll get the precise location and category of the shop, as well as details such as price and rating where necessary. Some of these calls will be necessarily premium calls, so to keep ourselves within the boundaries of free usage we'll limit them.

When using the search and explore endpoints the Foursquare API will naturally constrain the responses, presumably to keep people from scraping/farming locations. We'll need to keep this in mind, as it means that our results will not be "complete" for a given radius.

We'll begin by retrieving the full list of categories from Foursquare and identifying which of them corresponds to jewelry (note in the UK this is spelled "jewellery").

Out[7]:

	id	name
746	4bf58dd8d48988d111951735	Jewelry Store

We want to focus on a specific types of venues, and the Foursquare API allows us to break things down into category "sections": shops, food, drinks, coffee, arts. We'll make separate calls for each of these venue sections and collect the relevant details. Although this may not get us everything, narrowing the search on sections will get us significantly more options.

Then we'll grab all the Jewelry Shop category venues in London exclusively. Below is a sample of these stores, and a heatmap of all the stores across London.

Out[17]:

	Placename	Latitude	Longitude	Venue ID	Venue	Venue Latitude	Venu Longitude
0	Warren Street	51.5136	-0.1498	4ac518eff964a52063ad20e3	Grays Antiques	51.513622	-0.14875
1	Mayfair	51.5095	-0.1490	4ac518eff964a52066ad20e3	Montblanc Boutique	51.508787	-0.14084
2	Soho	51.5144	-0.1354	4ac518eff964a5208aad20e3	Storm	51.513183	-0.13876
3	Mayfair	51.5095	-0.1490	4ac518f0f964a520aaad20e3	Tiffany & Co.	51.509579	-0.14130
4	Oxford Circus	51.5154	-0.1414	4ac518f0f964a520c9ad20e3	H. Samuel	51.515108	-0.14325

Out[18]:



## 3. Methodology

As our interest is in finding the most interesting areas in London, we will use an unstructured approach, and after some initial exploratory work we will apply clustering strategies to the data, with the aim of uncovering patterns that we'll be able to use.

### 3.1. Exploratory Analysis

There are some basic principles of being a hipster, that we should be able to reflect in our preparation.

1. Chain stores are bad. Be they coffee shops, book stores or otherwise, if it's a major brand (with few exceptions) then they should be avoided.
2. Funky independent shops are good. Particularly those of a specific type, such as record stores or organic markets.
3. Variety is important. The more options for shopping, food and drink, the better.
4. Don't forget culture. An area that has an interest beyond consumerism is cool, so arts and sights are key.

Let's begin by looking at the data and seeing whether we can extract features from it.

#### 3.1.1. Venue Frequency

As a possible metric for understanding whether a given venue is a chain, we'll look at the occurrence of the venue with the same name throughout London. We want to easily identify the difference between actual independent stores, stores that might have 2-3 franchises across the city, and massive chains. We want to severely penalise the latter.


We'll begin by looking at venues in the following Foursquare "sections": Shops, Coffee, Food, and Drinks.

##### **Shops**

Looking at the frequency distribution of all venues in the Shops group (focusing on incidence  $> 2$ ), we can see a fairly high incidence of chains. And if we look at the shops themselves we note that there are quite a few grocery store chains as well as other usual suspects. Note that there are some cases where a venue has a different category, but looking at the available samples these are acceptable categorizations and not an error.

Out[20]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x11ad8df60>



Out[21]:

Venue	Venue Categories	
Tesco Express	Grocery Store	307
Sainsbury's Local	Convenience Store	207
	Supermarket	207
	Grocery Store	207
Boots	Pharmacy	168
Co-op Food	Grocery Store	142
Paddy Power	Betting Shop	102
WHSmith	Bookstore	88
	Stationery Store	88
Carphone Warehouse	Mobile Phone Shop	80
Superdrug	Cosmetics Shop	79
	Pharmacy	79
Iceland	Grocery Store	75
Sainsbury's	Supermarket	72
	Grocery Store	72
Lidl	Discount Store	69
	Supermarket	69
M&S Simply Food	Grocery Store	62
	Convenience Store	62
	Food & Drink Shop	62

Name: count, dtype: int64




## Coffee

Next we'll look at venues in the Coffee group. These have a less dramatic frequency curve, but there are still quite a few chains of reasonable size. The ubiquitous Starbucks is present, although not at the top of the list.

Out[22]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x11d198a20>



Out[23]:

Venue	Venue Categories	
Costa Coffee	Coffee Shop	345
Starbucks	Coffee Shop	194
Pret A Manger	Coffee Shop	160
	Sandwich Place	160
Caffè Nero	Coffee Shop	153
JOE & THE JUICE	Juice Bar	23
Patisserie Valerie	Café	21
Black Sheep Coffee	Coffee Shop	21
Paul	Café	20
	Bakery	20
	Sandwich Place	20
Benugo	Sandwich Place	19
	Bakery	19
	Food Truck	19
	Deli / Bodega	19
	Coffee Shop	19
	Café	19
GAIL's Bakery	Bakery	19
Wild Bean Cafe	Café	18
Krispy Kreme	Donut Shop	16


Name: count, dtype: int64

## Food

Third up is the Food group. Also a significant number of chains showing up.

Out[24]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x11a815cc0>



Out[25]:

Venue	Venue Categories	
Pret A Manger	Sandwich Place	140
Subway	Sandwich Place	132
McDonald's	Fast Food Restaurant	100
KFC	Fried Chicken Joint	98
	Fast Food Restaurant	98
Greggs	Bakery	94
Nando's	Portuguese Restaurant	93
PizzaExpress	Pizza Place	82
Domino's Pizza	Pizza Place	82
GAIL's Bakery	Bakery	37
itsu	Sushi Restaurant	37
Franco Manca	Pizza Place	35
Pizza Hut	Pizza Place	33
wagamama	Asian Restaurant	33
Leon	Fast Food Restaurant	32
	Restaurant	32
	Café	32
Byron	Burger Joint	30
Paul	Bakery	29
	Café	29

Name: count, dtype: int64

## Drinks


The Drinks group is a stranger outcome, as there are not as many clear chains. Doing a little research reveals that **All Bar One**, **Slug and Lettuce** and **Executive Lounge Hilton** are definitely part of a chain, but many of the others have the same name as a coincidence. Or due to the lack of creativity in pub names.

Consequently we won't be able to use this approach to easily identify whether a bar is a chain, so we'll need to exclude this feature from the list.



Out[26]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x11d47ad30>



Out[27]:

Venue	Venue Categories	
All Bar One	Bar	17
The Crown	Pub	15
	Gastropub	15
Slug & Lettuce	Bar	13
Be At One	Bar	12
	Cocktail Bar	12
The White Hart	Gastropub	11
	Pub	11
	Hotel	11
The Red Lion	Pub	10
The George	Bar	10
	Pub	10
Simmons Bar	Bar	9
	Cocktail Bar	9
Executive Lounge Hilton	Hotel Bar	9
The Royal Oak	Gastropub	9
	Pub	9
Simmons Bar	Pub	9
The Victoria	Pub	9
	Gastropub	9

Name: count, dtype: int64

### 3.1.2. Hipster Factor

We'll establish some basic criteria for what qualifies as a hipster shop. It will need to be one of a selected set of categories listed below, and it will need to have a 'Venue Frequency' less than 10 - this will exclude obvious chain stores (which can occur with Bookstores and Thift / Vintage Stores, while maintaining flexibility for independents of a certain size.

Out[30]:

```
['Bookstore',  
'Thrift / Vintage Store',  
'Market',  
'Farmers Market',  
'Arts & Crafts Store',  
'Flea Market',  
'Record Shop',  
'Street Food Gathering',  
'Beer Store',  
'Organic Grocery',  
'Antique Shop']
```

Out[31]:

	Latitude	Longitude	Placename	Number of Hip Venues
178	51.5251	-0.0769	Shoreditch	32
123	51.5085	-0.1257	London	25
140	51.5134	-0.1234	Embankment	24
124	51.5085	-0.1249	Charing Cross	24
147	51.5142	-0.1247	Covent Garden	24
160	51.5175	-0.1204	Holborn	23
215	51.5431	-0.1499	Chalk Farm	22
168	51.5225	-0.0857	Bank	21
125	51.5095	-0.1958	Notting Hill Gate	20
210	51.5406	-0.1433	Camden Town	20

### 3.1.3. Culture

Counting the number of venues by area that are "arts" or "sights" venues will give us a good measure of the cultural options.

Sample of "arts":

Out[ 32 ]:

	Latitude	Longitude	Placename	Number of Art Venues
219	51.5134	-0.1234	Embankment	161
228	51.5144	-0.1354	Soho	157
200	51.5085	-0.1249	Charing Cross	154
199	51.5085	-0.1257	London	154
226	51.5142	-0.1247	Covent Garden	153
237	51.5166	-0.1308	Tottenham Court Road	146
231	51.5154	-0.1414	Oxford Circus	111
242	51.5175	-0.1204	Holborn	110
230	51.5148	-0.1461	Oxford Street	85
269	51.5268	-0.1333	Euston	79

## 3.2. Clustering

Our features of interest then should are:

1. number of shops
2. number of coffee venues
3. number of food venues
4. uniqueness of shops
5. uniqueness of coffee venues
6. uniqueness of food venues
7. number of hip venues
8. uniqueness of hip venues
9. number of drinks venues
10. number of arts venues
11. number of sights

As a metric to reflect the uniqueness of a venue, we'll take the 'Venue Frequency' and square it, and take the average of this across all the venues within a given location. We'll call this the Mean Frequency Squared (MFS for short). A lower value means the shops in the location are more unique / independent, whereas a higher value will mean the shops are more likely to be chains.

We will then apply k-means clustering to the locations and their feature data, to discern patterns within the data and to identify locations that are most like each other. Once this task is completed, we'll review the individual clusters and determine which is the cluster most interesting to us to open a shop, and look at possible next steps.

Our data (in advance of normalization) looks like this:


Out[33]:

	Latitude	Longitude	Shops Count	Shops MFS	Coffee Count	Coffee MFS	Food Count	Food MFS	Drinks Count	Art Cour
0	51.3133	0.0343	7	17111.428571	4.0	29757.0	4.0	1.000000	2.0	2.
1	51.3148	-0.1570	3	1.000000	0.0	0.0	4.0	1.750000	2.0	2.
2	51.3192	0.0712	4	1.000000	0.0	0.0	0.0	0.000000	1.0	1.
3	51.3200	-0.1409	12	10634.166667	5.0	4682.6	11.0	6.363636	4.0	4.
4	51.3264	-0.1011	2	10082.500000	0.0	0.0	10.0	1.000000	2.0	4.

When we run k-means clustering against the data, we trial several different cluster numbers and review the resulting metrics to choose a good value for k.

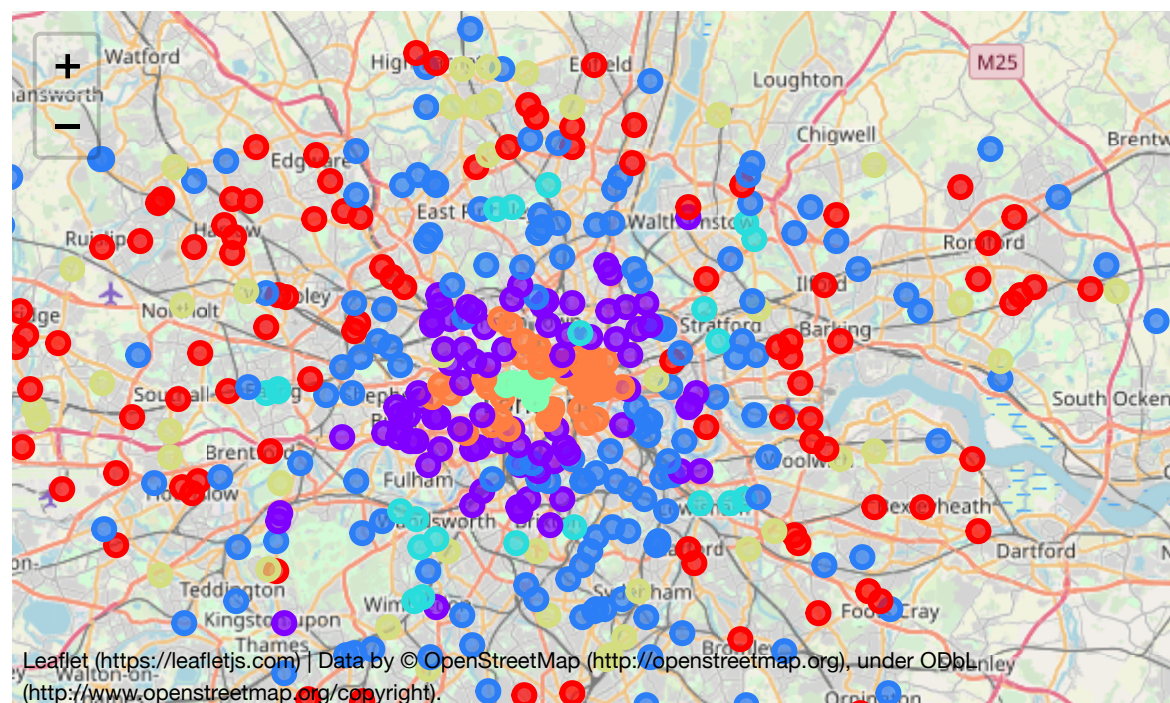
As we review the sum of square differences, there's not a sufficiently clear "elbow" to decide the value of k. Instead we use the silhouette score as a different metric and this gives us a bit more clarity.

Remarkably 2 metrics for k=5 and k=7 are quite close to each other, and we choose the slightly higher one rather than the one with the best score - the difference is marginal and the outcome is more finely tuned to our purposes.



We then run the clustering for the value of k=7, and produce a map of London overlaid with different locations color-coded with their cluster. Below follows a review of the individual clusters and their prospective value to us.

Out[39]:



## Cluster 0

Our first cluster is the second worst candidate to open a shop. It's the second least hip, and although it has more shops and these are slightly more independent, the food and coffee options are the most likely to be chains. This is likely to be a suburban area with slightly more activity, but an unwise place to open a hip store.

```
### cluster number: 0 has 98 locations:
```

```
Shops Count      25.489796
Shops MFS        6196.007732
Coffee Count     7.122449
Coffee MFS       34730.426611
Food Count       17.336735
Food MFS         2333.200831
Drinks Count     6.234694
Arts Count       2.765306
Sights Count     30.428571
Hipster Count    0.520408
Hipster MFS      0.945578
dtype: float64
```

```
# sample locations:
```

```
5      Sanderstead
7      Purley
20     South Beddington
25     Orpington
28     Worcester Park
29     Tolworth
30     Sutton
31     Croydon
45     Bromley
55     Chislehurst
Name: Placename, dtype: object
```

## Cluster 1

This cluster is the most likely candidate for us to open a shop. It has a reasonable cultural profile, a good number of coffee, food places and shops with independence. It also has decent number of hipster options with a good score.

```
### cluster number: 1 has 82 locations:
```

```
Shops Count      70.195122
Shops MFS        7039.706114
Coffee Count     45.902439
Coffee MFS       9282.302691
Food Count       98.914634
Food MFS         865.451266
Drinks Count     43.634146
Arts Count       13.865854
Sights Count     142.012195
Hipster Count    6.951220
Hipster MFS      4.253009
dtype: float64
```

```
# sample locations:
```

```
49      Kingston upon Thames
60           South Wimbledon
103           Tulse Hill
105      Richmond upon Thames
107           Richmond
110           Lambeth
111           Clapham
114      Clapham Junction
117           Brixton
118      Clapham North
Name: Placename, dtype: object
```

## Cluster 2

This cluster is one of the weaker candidate for opening a shop. It has few shops, coffee and food options, and culturally it's poor compared to some of the stronger clusters.

```
### cluster number: 2 has 162 locations:
```

```
Shops Count      19.882716
Shops MFS        6009.353817
Coffee Count     9.049383
Coffee MFS       8100.492138
Food Count       21.129630
Food MFS         537.102504
Drinks Count     8.425926
Arts Count       3.283951
Sights Count     34.839506
Hipster Count    1.296296
Hipster MFS      1.149544
dtype: float64
```

```
# sample locations:
```

```
1      Hooley
2      Cudham
3      Coulsdon
4      Kenley
6      Downe
8      Sutton
10     Addington
13     Chelsfield
14     Farnborough
15     Keston
Name: Placename, dtype: object
```

### Cluster 3

This cluster is a step up from our three weakest. It has a small amount of hip and cultural activity, a reasonable number of shops with decent independence, but it's the second most likely to have coffee and food options that are chains. Based purely on our metrics this is a slightly more urbanized suburbia, and it's telling that the weakest clusters for our purposes are all around the edge of London.



```
### cluster number: 3 has 22 locations:
```

```
Shops Count      57.636364
Shops MFS        5028.117808
Coffee Count     21.772727
Coffee MFS       13272.275323
Food Count       49.136364
Food MFS         1140.840503
Drinks Count     17.454545
Arts Count       5.772727
Sights Count     74.500000
Hipster Count    4.681818
Hipster MFS      30.447658
dtype: float64
```

```
# sample locations:
```

```
61      Wimbledon
66      Merton
84      Southfields
88      Wandsworth Common
93      West Hill
101     North Dulwich
106     Wandsworth
108     Putney
115     Lewisham
116     Blackheath
Name: Placename, dtype: object
```

#### Cluster 4

This is the beating heart of cultural activity and hipness in London, by our metrics - although it does lose a slight amount on the fact that the hipness has some chain stores. It has a massive number of shops, coffee and food options, with a mix of independent and chain options. Because of the large number of shops however the competition (also rent and business rates) would be very high, making this not a preferred location for our shop opening.

```
### cluster number: 4 has 9 locations:
```

```
Shops Count      191.000000
Shops MFS        2995.533499
Coffee Count     190.555556
Coffee MFS       5021.251725
Food Count       249.111111
Food MFS         323.768650
Drinks Count     227.444444
Arts Count       136.777778
Sights Count     226.222222
Hipster Count    19.777778
Hipster MFS      7.589503
dtype: float64
```

```
# sample locations:
```

```
215          London
216      Charing Cross
235      Embankment
242      Covent Garden
244          Soho
246      Oxford Street
248      Oxford Circus
254  Tottenham Court Road
259          Holborn
Name: Placename, dtype: object
```

## Cluster 5

This cluster is the worst candidate for us to open a shop. It's the least hip, has no activity to draw any kind of footfall. The few existing shops are all chains, although the coffee and food venues are more independent. Looking at the map representation, they are almost entirely outside of central London, so this is not surprising - the shops will more likely than not be grocery stores, and restaurants will be locally owned. This is also the least culturally active location.

```
### cluster number: 5 has 48 locations:
```

```
Shops Count      10.187500
Shops MFS        20014.968975
Coffee Count     5.270833
Coffee MFS       8974.225721
Food Count      11.520833
Food MFS        961.352942
Drinks Count     5.104167
Arts Count       2.312500
Sights Count    19.937500
Hipster Count    0.187500
Hipster MFS      0.187500
dtype: float64
```

```
# sample locations:
```

```
0      Biggin Hill
9      Wallington
11     Belmont
12     Cheam
16     Chessington
17     Malden Rushett
21     Hook
35     Morden Hall Park
36     Woodside
40     South Norwood
Name: Placename, dtype: object
```

## Cluster 6

This is one of the better clusters on our metrics count. It has a lot of food and drink options, and the former are quite independent. The coffee options are also significant, although there are quite a lot of chains present. The shops are significant but they are the second worst among all the clusters for independence. If you were looking to open a shop that was not so focused on uniqueness, you would consider this cluster as a viable option.

```
### cluster number: 6 has 47 locations:
```

```
Shops Count      126.085106
Shops MFS        8102.034243
Coffee Count     122.510638
Coffee MFS       10992.274526
Food Count       205.191489
Food MFS         730.231981
Drinks Count     138.808511
Arts Count       47.212766
Sights Count     221.319149
Hipster Count    12.659574
Hipster MFS      5.348187
dtype: float64
```

```
# sample locations:
```

```
179      South Kensington
183              Victoria
188              London
190  Westminster Bridge
194      Marylebone
195      Westminster
196      Southwark
206      South Bank
207      Southwark
209      Waterloo
Name: Placename, dtype: object
```

## 4. Results

Based on our evaluation, Cluster 1 is the most likely candidate for a place to open a shop.

It's worth refining this a bit more to break down where to start our search.

### The "Low Competition" Approach

One possible way to look at this is to see whether there are other Jewelry shops in the area - look for places with less or no competition but which still have a good amount of other cool shops.

Given that approach, our top 5 options are:

Out[54]:

	Latitude	Longitude	Placename	Postalcode	Hipster Count	Hipster MFS	Jewel/Shop Density
338	51.5480	-0.0629	Shacklewell	E8	18.0	7.111111	0.0
288	51.5272	-0.0611	Bethnal Green	E2	15.0	2.066667	0.0
214	51.5075	-0.2050	Ladbroke Grove	W11	13.0	4.307692	0.0
173	51.4911	-0.1493	Belgravia	SW1	12.0	1.000000	0.0
103	51.4548	-0.1158	Tulse Hill	SW2	11.0	1.000000	0.0

### The "High Footfall" Approach

One other approach might be to worry less about the competition and focus on the total number of venues in the area. This strategy would assume that your goal is footfall, implying that more things to do in the area means more people coming your way.

Given that approach, our top 5 options are:

Out[56]:

	Latitude	Longitude	Placename	Postalcode	Total Count
49	51.4126	-0.2974	Kingston upon Thames	KT1	286.0
312	51.5362	-0.1030	Islington	EC1P,EC4P,EC50,N1,N1C,N1P,N7	275.0
210	51.5050	-0.2211	Western Avenue	W12	241.0
202	51.5015	-0.1619	Knightsbridge	SW1	223.0
162	51.4870	-0.1910	West Brompton	SW5	219.0

### The "Super Hipster" Approach

One final approach might be to just aim for the hippest place you can find, and not worry about anything else. So focus on the number of hip shops and their independence.

Given that approach, our top 5 options are:

Out[57]:

	Latitude	Longitude	Placename	Postalcode	Hipster Count	Hipster MFS	Jewel/Shop Density
309	51.5344	-0.0694	Shoreditch	N1	18.0	1.166667	0.072464
338	51.5480	-0.0629	Shacklewell	E8	18.0	7.111111	0.000000
226	51.5115	-0.2059	Notting Hill	W11	16.0	3.687500	0.027027
288	51.5272	-0.0611	Bethnal Green	E2	15.0	2.066667	0.000000
292	51.5279	-0.1072	Finsbury	EC1	14.0	8.857143	0.011364

## 5. Discussion

This evaluation was done within the constraints of the Foursquare API and the level of free access available. A more complete/mature approach would be to retrieve ratings and price for venues, and to use that to identify the affluence of the specific neighbourhood based on that. We might find that even though an area is moderately populated with shops, these are much more high end. And conversely we may find that an area that has a large number of shops with a high amount of independence is because they are in a poorer area, and therefore less likely to be interested in buying jewellery.

The Foursquare API also had a limited set returned for venues in a particular area, which was not always consistent. Radial searches might return the same number of results even though the radius was increased. There was also a fair amount of overlap on locations, as well as some areas on the map that had limited coverage. More time could be spent on selecting the locations to ensure full map coverage, but also to align with consumer behavior such as focusing on a known set of shopping areas.

We could also retrieve demographic data from a different source such as the UK Census. At the time of this evaluation the most recent data was for 2011 and therefore deemed to be too old to be relevant, but new data will be released within the next 2 years and will be much more up-to-date.

There may also be more useful econographic data that could inform a stronger model. This data often requires a fee, but a student license might be available for the purposes of this study.

Commercial and residential property prices by area would be an incredibly powerful tool, as it would allow for a clear filter of what would be an acceptable price to pay and what a target turnover for a shop would need to be if it opened.

Finally, further exploratory work could be done to build useful features, with the data we have available. A more careful study of the categories available in the given larger groups might mean we exclude categories that are not interesting or useful to the study, or score specific categories higher than others based on their importance.

## 6. Conclusion

We were able to successfully build some metrics of interest to an independent jewelry shop, and use k-means clustering to identify locations in Greater London that are good targets for this kind of business.

Ironically, the original first location for the shop is in 'Shoreditch', which was rated the hippest location in Cluster 1.

So then... Data Science is hip.